

Statistical Machine Learning (BE4M33SSU)

Hidden Markov Models

Czech Technical University in Prague: Boris Flach and Jan Drchal

- ◆ Markov models on sequences: inference & parameter learning
- ◆ Hidden Markov Models (HMMs)

Structured Hidden States

Models discussed so far: mainly classifiers predicting a categorical (class) variable $y \in \mathcal{Y}$

Often in applications: the hidden state y is a structured variable.

Here: the hidden state y is a **sequence** of categorical variables.

Application examples:

- ◆ language modelling (next token prediction),
- ◆ text recognition (printed, handwritten),
- ◆ token classification (named entity recognition, POS tagging),
- ◆ speech recognition (single word recognition, continuous speech recognition, translation),
- ◆ text-to-speech,
- ◆ bioinformatics (gene prediction),
- ◆ finance (predict stock price),
- ◆ cybersecurity (anomaly detection),
- ◆ robot self-localization.

Markov Models and Hidden Markov Models on chains:

a class of generative probabilistic models for sequences of features and sequences of categorical variables.

Markov Models: Preliminaries

Let $\mathbf{s} = (s_1, s_2, \dots, s_n)$ denote a sequence of length n with elements from a finite set K .

Each element corresponds to a *discrete time step*.

Any joint probability distribution on K^n can be written as

$$p(s_1, s_2, \dots, s_n) = p(s_1) p(s_2 | s_1) p(s_3 | s_2, s_1) \cdot \dots \cdot p(s_n | s_1, \dots, s_{n-1})$$

To see that, use the definition of conditional probability recursively:


$$\begin{aligned} p(s_1, s_2, \dots, s_n) &= p(s_n | s_1, \dots, s_{n-1}) p(s_1, s_2, \dots, s_{n-1}) = \\ &= p(s_n | s_1, \dots, s_{n-1}) p(s_{n-1} | s_1, s_2, \dots, s_{n-2}) p(s_1, s_2, \dots, s_{n-2}) = \dots \end{aligned}$$

Markov Models: Definition

Definition 1. A joint p.d. on K^n is a *first-order* Markov model if

$$p(\mathbf{s}) = p(s_1) p(s_2 | s_1) p(s_3 | s_2) \cdot \dots \cdot p(s_n | s_{n-1}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$

holds for any $\mathbf{s} = (s_1, s_2, \dots, s_n)$.




Remark 1. Higher degree MMs consider a longer history, e.g., a *second-order* MM uses $p(s_t | s_{t-1}, s_{t-2})$ *transition probabilities*.

Markov Models: Definition

Definition 1. A joint p.d. on K^n is a *first-order* Markov model if

$$p(\mathbf{s}) = p(s_1) p(s_2 | s_1) p(s_3 | s_2) \cdot \dots \cdot p(s_n | s_{n-1}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$

holds for any $\mathbf{s} = (s_1, s_2, \dots, s_n)$.



Remark 1. Higher degree MMs consider a longer history, e.g., a *second-order* MM uses $p(s_t | s_{t-1}, s_{t-2})$ *transition probabilities*.


Question 1: How to generate a sequence?

Markov Models: Definition

Definition 1. A joint p.d. on K^n is a *first-order* Markov model if

$$p(\mathbf{s}) = p(s_1) p(s_2 | s_1) p(s_3 | s_2) \cdot \dots \cdot p(s_n | s_{n-1}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$

holds for any $\mathbf{s} = (s_1, s_2, \dots, s_n)$.



Remark 1. Higher degree MMs consider a longer history, e.g., a *second-order* MM uses $p(s_t | s_{t-1}, s_{t-2})$ *transition probabilities*.

Question 1: How to generate a sequence?


Question 2: What is the number of sequences for English alphabet?

Markov Models: Definition

Definition 1. A joint p.d. on K^n is a *first-order* Markov model if

$$p(\mathbf{s}) = p(s_1) p(s_2 | s_1) p(s_3 | s_2) \cdot \dots \cdot p(s_n | s_{n-1}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$

holds for any $\mathbf{s} = (s_1, s_2, \dots, s_n)$.



Remark 1. Higher degree MMs consider a longer history, e.g., a *second-order* MM uses $p(s_t | s_{t-1}, s_{t-2})$ *transition probabilities*.

Question 1: How to generate a sequence?


Question 2: What is the number of sequences for English alphabet?
 26^n for 26 lowercase letters.

Markov Models: Definition

Definition 1. A joint p.d. on K^n is a *first-order* Markov model if

$$p(\mathbf{s}) = p(s_1) p(s_2 | s_1) p(s_3 | s_2) \cdot \dots \cdot p(s_n | s_{n-1}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$

holds for any $\mathbf{s} = (s_1, s_2, \dots, s_n)$.



Remark 1. Higher degree MMs consider a longer history, e.g., a *second-order* MM uses $p(s_t | s_{t-1}, s_{t-2})$ *transition probabilities*.

Question 1: How to generate a sequence?

Question 2: What is the number of sequences for English alphabet?

26^n for 26 lowercase letters.


Question 3: How many values do we need to fix the *transition probabilities* $p(s_t | s_{t-1})$?

Markov Models: Definition

Definition 1. A joint p.d. on K^n is a *first-order* Markov model if

$$p(\mathbf{s}) = p(s_1) p(s_2 | s_1) p(s_3 | s_2) \cdot \dots \cdot p(s_n | s_{n-1}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$

holds for any $\mathbf{s} = (s_1, s_2, \dots, s_n)$.



Remark 1. Higher degree MMs consider a longer history, e.g., a *second-order* MM uses $p(s_t | s_{t-1}, s_{t-2})$ *transition probabilities*.

Question 1: How to generate a sequence?

Question 2: What is the number of sequences for English alphabet?

26^n for 26 lowercase letters.

Question 3: How many values do we need to fix the *transition probabilities* $p(s_t | s_{t-1})$?
 $K \times K$ (*stochastic matrix*: columns sum to 1) $\Rightarrow K \times (K - 1)$ parameters is enough.

Markov Models: Matrix Representation

The elements of a **state distribution** vector π_t for time step t are defined as:

$$(\pi_t)_j = p(s_t = j).$$

Similarly, a **transition matrix** P is defined as:

$$P_{ij}(t) = p(s_t = i \mid s_{t-1} = j),$$

where $\sum_i P_{ij}(t) = 1$, i.e., the P is a *stochastic matrix*.

A Markov model is called **homogeneous** if the transition probabilities P do not depend on time. In this case, the *state distribution* at time step t can be easily written as:

$$\pi_t = P^{t-1} \pi_1,$$

where π_1 is the *initial state (starting) distribution*.

Markov Models: Random Walker Example

Example 1 (Random walk on a graph).

- ◆ Let (V, E) be a directed graph. A random walk in (V, E) is described by a sequence $s = (s_1, \dots, s_t, \dots)$ of visited nodes, i.e., $s_t \in V$.
- ◆ The walker starts in node $i \in V$ with probability $p(s_1 = i)$.
- ◆ The edges of the graph are weighted by $w : E \rightarrow \mathbb{R}_+$, s.t.

$$\sum_{j: (i,j) \in E} w_{ij} = 1 \quad \forall i \in V$$

- ◆ In the current position $s_t = i$, the walker randomly chooses an outgoing edge with probability given by the weights and moves along this edge, i.e.

$$p(s_{t+1} = j \mid s_t = i) = \begin{cases} w_{ij} & \text{if } (i, j) \in E \\ 0 & \text{otherwise} \end{cases}$$




Questions: How does the distribution $p(s_t)$ behave? Does it converge to some fix-point distribution π for $t \rightarrow \infty$?

Return Times and Limiting Distributions

- ◆ A homogeneous Markov model is **irreducible** if each state l can be reached starting from any state k with non-zero probability (after some number of transitions).
- ◆ A state k has **return time** τ if it can be reached with non-zero probability after τ transitions when starting from itself.
- ◆ A state $k \in K$ is **a-periodic** if the greatest common divisor of its *return times* is 1.

Example 2. Markov chains:

- ◆ **left:** *irreducible, a-periodic* chain: $d(1) = d(2) = \gcd(2, 3, 4, 5, 6, \dots) = 1$, $d(3) = \gcd(3, 5, 6, \dots) = 1$.
- ◆ **middle:** *reducible* chain: starting in 4 (*absorbing state*) we end up with $\pi = (0, 0, 0, 1)$; starting in 1 or 2 we get $\pi = (\frac{1}{2}, \frac{1}{2}, 0, 0)$; starting in 3 leads to either of the two *stationary distributions*.
- ◆ **right:** *irreducible* chain for $\alpha, \beta = 0.9$: we end up with $\pi = (\frac{1}{2}, \frac{1}{2})$ by symmetry. For $\alpha, \beta = 1$ it depends where you start.



Return Times and Limiting Distributions, contd.

8/29

Theorem 1. Let P be the transition probability matrix of an **irreducible** homogeneous Markov model with all states **a-periodic**. Then there exists a unique **stationary distribution** π^* s.t. $P\pi^* = \pi^*$. Moreover, it is a **limiting distribution**, i.e.

$$\lim_{t \rightarrow \infty} P^t \pi_1 = \pi^*$$

for arbitrary starting distributions π_1 .

Remark 2. The two assumptions for the limiting distribution:

- ◆ the **irreducibility** means that the state transition diagram must be a singly-connected component,
- ◆ the **a-periodicity** means no oscillations.

Remark 3. A special case of the above is a **regular** chain for which: $P_{ij}^n > 0$ for some n and all i and j :

- ◆ Hence, after n steps, we can end up in any state, no matter where we started.
- ◆ *Irreducibility* and a *self-loop* for all states is all we need.

Google PageRank


- ◆ One of the methods used to rank the retrieved documents (big in 1997).
- ◆ **Idea**: good pages are referenced more often; a random surfer spends more time on highly reachable pages.
- ◆ It is the *Random walk*:
 - P – transition matrix based on hyperlinks between pages,
 - π_i – page i authority (high for being linked by many other pages),
 - authority is passed to the linked pages:

$$\pi_i = \sum_j P_{ij} \pi_j.$$

- Corresponds to the Markov model *stationary distribution*.
- ◆ We need the P to be *irreducible* having all states *a-periodic*:
 - Real link adjacency matrix G is sparse as not all pages are linked to all other pages.
 - We must set P , so even non-existing links (in G) get at least a very small probability, so we end up with a **regular** chain.



Google PageRank Example

Simple example



Converges to $\pi = (0.3209, 0.1706, 0.1065, 0.1368, 0.0643, 0.2008)$

Example for 500 pages linked from www.harvard.edu:
Adjacency matrix G (left), PageRank vector π (right).



Algorithms: Computing the Most Probable Sequence


We want to compute the most probable sequence $\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} \left[p(s_1) \prod_{t=2}^n p(s_t | s_{t-1}) \right]$.

Take the logarithm of $p(\mathbf{s})$: $\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} \left[g_1(s_1) + \sum_{t=2}^n g_t(s_{t-1}, s_t) \right]$

and apply dynamic programming: Set $\phi_1(s_1) \equiv g_1(s_1)$ and compute

$$\phi_t(s_t) = \max_{s_{t-1} \in K} \left[\phi_{t-1}(s_{t-1}) + g_t(s_{t-1}, s_t) \right] \quad \forall s_t \in K.$$

Finally, find $s_n^* \in \arg \max_{s_n \in K} \phi_n(s_n)$ and back-track the solution. This corresponds to searching the best path in the graph



Question: What is the run-time complexity? .

Algorithms: Computing the Most Probable Sequence


We want to compute the most probable sequence $\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} \left[p(s_1) \prod_{t=2}^n p(s_t | s_{t-1}) \right]$.

Take the logarithm of $p(\mathbf{s})$: $\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} \left[g_1(s_1) + \sum_{t=2}^n g_t(s_{t-1}, s_t) \right]$

and apply dynamic programming: Set $\phi_1(s_1) \equiv g_1(s_1)$ and compute

$$\phi_t(s_t) = \max_{s_{t-1} \in K} \left[\phi_{t-1}(s_{t-1}) + g_t(s_{t-1}, s_t) \right] \quad \forall s_t \in K.$$

Finally, find $s_n^* \in \arg \max_{s_n \in K} \phi_n(s_n)$ and back-track the solution. This corresponds to searching the best path in the graph




Question: What is the run-time complexity? $\mathcal{O}(nK^2)$.

Algorithms: Computing Marginal Probabilities

How to compute marginal probabilities for the sequence element s_j in position j

$$p(s_j) = \sum_{s_1 \in K} \cdots \cancel{\sum_{s_j \in K}} \cdots \sum_{s_n \in K} p(s_1) \prod_{t=2}^n p(s_t | s_{t-1})$$



Summation over the **trailing** variables is easily done because:

$$\sum_{s_n \in K} p(s_1) \cdots p(s_{n-1} | s_{n-2}) p(s_n | s_{n-1}) = p(s_1) \cdots p(s_{n-1} | s_{n-2})$$

Algorithms: Computing Marginal Probabilities, contd.

The summation over the **leading** variables is done dynamically: Begin with $p(s_1)$ and compute:

$$p(s_t) = \sum_{s_{t-1} \in K} p(s_t | s_{t-1}) p(s_{t-1}) \quad \forall s_t \in K.$$

Equivalently using matrix representation the computation reads: $\pi_t = P(t)\pi_{t-1}$.

Run-time complexity once again $\mathcal{O}(nK^2)$.

Remark 4.

- ◆ Notice that the preferred direction (from first to last) in the definition of a Markov model is only apparent. By computing the marginal probabilities $p(s_t)$ and by using $p(s_t | s_{t-1})p(s_{t-1}) = p(s_{t-1}, s_t) = p(s_{t-1} | s_t)p(s_t)$, we can rewrite the model in reverse order.

Algorithms: Learning a Markov Model

Suppose we are given i.i.d. training data $\mathcal{T}^m = \{\mathbf{s}^k \in K^n \mid k = 1, \dots, m\}$ and want to estimate the parameters of the Markov model by the maximum likelihood estimate. This is very easy:

- ◆ Denote by $a(s_{t-1} = j, s_t = i)$ the number of sequences in $\mathcal{T}^m$ for which $s_{t-1} = j$ and $s_t = i$.
- ◆ The estimates for the conditional probabilities are then given by

$$p(s_t = i \mid s_{t-1} = j) = \frac{a(s_{t-1} = j, s_t = i)}{\sum_k a(s_{t-1} = j, s_t = k)}.$$

Algorithms: Learning a Markov Model, contd.

But is our estimate an MLE?

Let's recall the log-likelihood:

$$\frac{1}{m} \sum_{s \in \mathcal{T}^m} \left[\log p(s_1) + \sum_{t=2}^n \log p(s_t | s_{t-1}) \right]$$

Proof (idea):

Consider all terms in the log-likelihood that depend on the transition probability from $(t-1) \rightarrow t$ and rewrite them using transition counts $a(s_{t-1} = j, s_t = i)$

$$\frac{1}{m} \sum_{s \in \mathcal{T}^m} \log p(s_t | s_{t-1}) = \frac{1}{m} \sum_{i,j \in K} a(s_{t-1} = j, s_t = i) \log p(s_t = i | s_{t-1} = j)$$

Maximise this w.r.t. $p(s_t | s_{t-1})$ under the constraint $\sum_{s_t \in K} p(s_t | s_{t-1}) = 1$ using Lagrange multipliers.

Markov Models are Exponential Families

Markov models are **exponential families**. For simplicity we show this for the family of homogeneous Markov models on sequences $\mathbf{s} = (s_1, s_2, \dots, s_n)$ of length n under the additional assumption that $p(s_1) = \frac{1}{K}$.

We have

$$p(\mathbf{s}) = \frac{1}{K} \prod_{t=2}^n p(s_t | s_{t-1}),$$

- ◆ sufficient statistic: $\Phi(\mathbf{s})$ is a $K \times K$ matrix with entries $\Phi_{kl}(\mathbf{s})$ counting the number of transitions from state l to state k in the sequence $\mathbf{s}$,
- ◆ natural parameter: H is a $K \times K$ matrix with entries $H_{kl} = \log p(s_t = k | s_{t-1} = l)$,
- ◆ base measure: $h(\mathbf{s}) = 1$,
- ◆ cumulant function: $\log(K)$.

We can write the probability of sequences as

$$p(\mathbf{s}; H) = \exp[\langle \Phi(\mathbf{s}), H \rangle - \log(K)]$$

Remark 5. This can be generalised for models with non-uniform $p(s_1)$ and also for general (i.e. non-homogeneous) Markov models.

Hidden Markov Models

- ◆ Let $\mathbf{s} = (s_1, s_2, \dots, s_n)$ denote a sequence of hidden states from a finite set K .
- ◆ Let $\mathbf{x} = (x_1, x_2, \dots, x_n)$ denote a sequence of features from some feature space $\mathcal{X}$.

Definition 2. A joint p.d. on $\mathcal{X}^n \times K^n$ is a Hidden Markov model if

- (a) the prior p.d. $p(\mathbf{s})$ for the sequences of hidden states is a Markov model, and
- (b) the conditional distribution $p(\mathbf{x} | \mathbf{s})$ for the feature sequence is independent, i.e.


$$p(\mathbf{x} | \mathbf{s}) = \prod_{t=1}^n p(x_t | s_t).$$

The joint model is given by

$$p(\mathbf{x}, \mathbf{s}) = p(s_1) \prod_{t=2}^n p(s_t | s_{t-1}) \prod_{t=1}^n p(x_t | s_t).$$

Hidden Markov Model: OCR Example

Example 3 (Text recognition, OCR).




- ◆ $\mathbf{x} = (x_1, x_2, \dots, x_n)$ – sequence of images with characters,
- ◆ $\mathbf{s} = (s_1, s_2, \dots, s_n)$ – sequence of alphabetic characters,
- ◆ $p(s_t | s_{t-1})$ – language model (*transition probabilities*),
- ◆ $p(x_t | s_t)$ – appearance model for characters (*emission probabilities*).

Question: How to generate a sequence?

Hidden Markov: Inference Tasks

- ◆ **Filtering**: compute the **belief state** $p(s_t | \mathbf{x}_{1:t})$ *online* using **forward** algorithm. Get less noise than using just $p(s_t | x_t)$.
- ◆ **Smoothing**: compute $p(s_t | \mathbf{x}_{1:n})$ *offline* given all evidence using **forward-backward** algorithm. Even less uncertainty given future data.
- ◆ **Fixed lag smoothing**: $p(s_{t-l} | \mathbf{x}_{1:t})$, where $l > 0$ is called the lag, compromise of online and offline.
- ◆ **Most probable sequence**: compute $\arg \max_{\mathbf{s}_{1:n}} p(\mathbf{s}_{1:n} | \mathbf{x}_{1:n})$ using **Viterbi** algorithm.
- ◆ **Prediction**: compute $p(s_{t+h} | \mathbf{x}_{1:t})$, where $h > 0$ is called the prediction **horizon**: just use the *belief state* and power of the transition matrix P^h (for a homogeneous HMM).
- ◆ **Posterior samples**: generate new samples $\mathbf{s}_{1:n}^s \sim p(\mathbf{s}_{1:n} | \mathbf{x}_{1:n})$.



Filtering: Forward Algorithm

We want to compute $\alpha_t(k) \triangleq p(s_t = k | \mathbf{x}_{1:t})$, $\forall k \in K$ in some sequence position t .

Let us start with one-step-ahead predictive density:

$$p(s_t = k | \mathbf{x}_{1:t-1}) = \sum_j \underbrace{p(s_t = k | s_{t-1} = j)}_{\Psi(j,k)} \underbrace{p(s_{t-1} = j | \mathbf{x}_{1:t-1})}_{\alpha_{t-1}(j)},$$

now define the update step:

$$\alpha_t(k) = p(s_t = k | \mathbf{x}_{1:t}) = p(s_t = k | x_t, \mathbf{x}_{1:t-1}) = \frac{1}{Z_t} p(x_t | s_t = k, \mathbf{x}_{1:t-1}) p(s_t = k | \mathbf{x}_{1:t-1}),$$

where $Z_t = p(x_t | \mathbf{x}_{1:t-1}) = \sum_{k'} p(x_t | s_t = k') p(s_t = k' | \mathbf{x}_{1:t-1})$.

The matrix-vector form of the update is:

$$\alpha_t \propto \psi_t \odot (\Psi^T \alpha_{t-1}),$$

where $\psi_t(k) \triangleq p(x_t | s_t = k)$ denotes the emission probabilities and $\odot$ is the Hadamard product (elementwise vector multiplication). Initialize with $\alpha_1 \propto \psi_1 \odot \pi_1$.

This is the **forward** algorithm.

Smoothing: Forward-Backward Algorithm

How to compute marginal probabilities for hidden states given the sequence of features?

$$p(s_t = k | \mathbf{x}) = p(s_t = k | \mathbf{x}_{1:t}, \mathbf{x}_{t+1:n}) \propto \underbrace{p(s_t = k | \mathbf{x}_{1:t})}_{\text{forward}} \underbrace{p(\mathbf{x}_{t+1:n} | s_t = k, \mathbf{x}_{1:t})}_{\text{backward}}$$

We know how to compute the *forward*, now let us look at the *backward* part, firstly define the conditional likelihood of future evidence given that the hidden state at time i :

$$\beta_t(k) \triangleq p(\mathbf{x}_{t+1:n} | s_t = k).$$

Let us establish the recurrence:

$$\begin{aligned} \beta_{t-1}(k) &= p(\mathbf{x}_{t:n} | s_{t-1} = k) = \sum_j p(s_t = j, \mathbf{x}_t, \mathbf{x}_{t+1:n} | s_{t-1} = k) = \\ &= \sum_j p(\mathbf{x}_{t+1:n} | s_t = j, \cancel{s_{t-1} = k}, \cancel{\mathbf{x}_t}) p(s_t = j, \mathbf{x}_t | s_{t-1} = k) = \\ &= \sum_j p(\mathbf{x}_{t+1:n} | s_t = j) p(\mathbf{x}_t | s_t = j, \cancel{s_{t-1} = k}) p(s_t = j | s_{t-1} = k) = \sum_j \beta_t(j) \psi_t(j) \Psi(k, j). \end{aligned}$$

The **backward** algorithm is then $\beta_{t-1} \propto \Psi(\psi_t \odot \beta_t)$, with $\beta_n(k) = p(\underbrace{\mathbf{x}_{n+1:n}}_{\emptyset} | s_n = k) = 1$.

Smoothing: Forward-Backward Algorithm, contd.

The forward-backward algorithm now reads:

$$p(s_t = k | \mathbf{x}) \propto \alpha_t(k) \beta_t(k),$$

where both $\alpha_t(k)$ and $\beta_t(k)$ are computed using dynamic programming. The result must be normalized.

The computational complexity for computing $\alpha_t(k)$ and $\beta_t(k)$ in all positions $t = 1, \dots, n$ is thus $\mathcal{O}(nK^2)$.

Two-Slice Smoothed Marginals

We can also compute **pairwise** marginal probabilities $p(s_t, s_{t+1} | \mathbf{x})$ from the α -s and β -s. Start with:

$$\begin{aligned} p(s_t, s_{t+1}, \mathbf{x}) &= p(\mathbf{x}_{1:t} | s_t, s_{t+1}) p(x_{t+1} | s_t, s_{t+1}) p(\mathbf{x}_{t+2:n} | s_t, s_{t+1}) p(s_{t+1} | s_t) p(s_t) \\ &= p(s_t | \mathbf{x}_{1:t}) p(\mathbf{x}_{1:t}) p(x_{t+1} | s_{t+1}) p(\mathbf{x}_{t+2:n} | s_{t+1}) p(s_{t+1} | s_t) \end{aligned}$$

Now we get:

$$p(s_t = i, s_{t+1} = j | \mathbf{x}) = \frac{p(s_t = i, s_{t+1} = j, \mathbf{x})}{p(\mathbf{x})} \propto \alpha_t(i) p(x_{t+1} | s_{t+1}) p(s_{t+1} | s_t) \beta_{t+1}(j).$$

To summarize the computation decomposes to:

- ◆ $\alpha_t(i)$: past evidence,
- ◆ $p(s_{t+1} | s_t)$: transition to the *current* state,
- ◆ $p(x_{t+1} | s_{t+1})$: *current* observation,
- ◆ $\beta_{t+1}(j)$: future observations.

The computational complexity remains $\mathcal{O}(nK^2)$.

Most Probable Sequence: the Viterbi Algorithm

How to find the most probable sequence of hidden states given the sequence of features x

$$\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} p(s_1) \prod_{t=2}^n p(s_t | s_{t-1}) \prod_{t=1}^n p(x_t | s_t)$$

Take the logarithm, with

$$g_1(s_1, x_1) = \log[p(s_1)p(x_1 | s_1)]$$

$$g_t(s_{t-1}, s_t, x_t) = \log[p(s_t | s_{t-1})p(x_t | s_t)], \quad t = 2, \dots, n$$

Solve the task

$$\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} \left[g_1(s_1, x_1) + \sum_{t=2}^n g_t(s_{t-1}, s_t, x_t) \right]$$

by dynamic programming as before for Markov models.

The computational complexity is again $\mathcal{O}(nK^2)$.

Viterbi vs. Smoothing

Compare the most probable sequence computer using the Viterbi algorithm:

$$\mathbf{s}^* \in \arg \max_{\mathbf{s} \in K^n} p(s_1) \prod_{t=2}^n p(s_t | s_{t-1}) \prod_{t=1}^n p(x_t | s_t)$$

with the smoothed marginals computed using forward-backward:

$$\hat{\mathbf{s}} \triangleq \left(\arg \max_{s_1 \in K} p(s_1 | \mathbf{x}), \arg \max_{s_2 \in K} p(s_2 | \mathbf{x}), \dots, \arg \max_{s_n \in K} p(s_n | \mathbf{x}) \right).$$

They are not the same! Example with two time steps and observations = hidden states:

	$x_1 = 0$	$x_1 = 1$	
$x_2 = 0$	0.04	0.3	0.34
$x_2 = 1$	0.36	0.3	0.66
	0.4	0.6	

Viterbi: (0,1) is globally optimal, smoothed marginals (1,1) locally optimal.

Sampling Sequences from the Posterior

The task is to:

$$\mathbf{s}_{1:n}^s \sim p(\mathbf{s}_{1:n} \mid \mathbf{x}_{1:n})$$

We can follow these steps:

- ◆ Compute the *two-slice smoothed marginals* $p(s_t, s_{t+1} \mid \mathbf{x})$.
- ◆ Normalize them to get $p(s_{t+1} \mid s_t, \mathbf{x})$.
- ◆ Sample from the initial pair of states: $\mathbf{s}_{1:2}^s \sim p(s_1, s_2 \mid \mathbf{x})$.
- ◆ Recursively sample $s_{t+1}^s \sim p(s_{t+1} \mid s_t, \mathbf{x})$.

Note this is forwards-backwards pass and a additional sampling pass. It is possible to do it in a single pass...

Learning algorithms for HMMs

Supervised learning:

Given i.i.d. training data $\mathcal{T}^m = \{(\mathbf{x}^j, \mathbf{s}^j) \in \mathcal{X}^n \times K^n \mid j = 1, \dots, m\}$, estimate the parameters of the HMM by the maximum likelihood estimator.

This is done by simple “counting” as before for Markov models.

- ◆ Denote by $a_t(s_{t-1} = j, s_t = i)$ the number of examples in $\mathcal{T}^m$ for which $s_{t-1} = j$ and $s_t = i$.
- ◆ Denote by $b_t(s_t = i, x_t = x)$ the number of examples in $\mathcal{T}^m$ for which $s_t = i$ and $x_t = x$.

The estimates for the model parameters are then given by

$$p(s_t = i \mid s_{t-1} = j) = \frac{a_t(s_{t-1} = j, s_t = i)}{\sum_{i'} a_t(s_{t-1} = j, s_t = i')}, \quad p(x_t = x \mid s_t = i) = \frac{b_t(s_t = i, x_t = x)}{\sum_{x'} b_t(s_t = k, x_t = x')}$$

Remark 6. This is easy to generalise for the case that $\mathcal{X} = \mathbb{R}^m$ and $p(x_t \mid s_t)$ is from some parametric distribution family.

Learning algorithms for HMMs

Unsupervised learning:

Given i.i.d. training data $\mathcal{T}^m = \{\mathbf{x}^j \in \mathcal{X}^n \mid j = 1, \dots, m\}$, estimate the parameters of the HMM by the maximum likelihood estimator.

We apply the EM-algorithm (aka Baum-Welch algorithm): Initialise the model, then iterate

E-step Use the current model estimate and compute the two-slice smoothed marginals

$$\alpha_{\mathbf{x}}(s_t, s_{t-1}) = p(s_t, s_{t-1} \mid \mathbf{x}) \quad s_{t-1}, s_t \in K$$

for each training example $\mathbf{x} \in \mathcal{T}^m$ and all positions $i = 2, \dots, n$.

M-step Use the α -s as soft labels for computing the counts as in supervised learning

$$a_t(s_{t-1}, s_t) = \sum_{\mathbf{x} \in \mathcal{T}^m} \alpha_{\mathbf{x}}(s_t, s_{t-1}), \quad b_t(s_t, x) = \sum_{\mathbf{x} \in \mathcal{T}(x_t=x)} \alpha_{\mathbf{x}}(s_t)$$

where $\mathcal{T}(x_t = x)$ is the set of training examples with $x_t = x$ and $\alpha_{\mathbf{x}}(s_t)$ is given by $\alpha_{\mathbf{x}}(s_t) = \sum_{s_{t-1}} \alpha_{\mathbf{x}}(s_t, s_{t-1})$.

Summary

- ◆ Hidden Markov models are statistical models for pairs of processes (sequences) (s, x) , where s is a sequence of hidden states and not directly observable.
- ◆ Markov models and HMMs are exponential families.
- ◆ Their parameters can be estimated by supervised learning using either Maximum likelihood estimates or Empirical risk minimization.
- ◆ Their parameters can be estimated by unsupervised learning using the EM-algorithm.
- ◆ All important inference and learning algorithms for HMMs have time complexity linear in the length of the sequence and quadratic in the size of the hidden state space.
- ◆ What if the space of hidden states is very large, e.g. $s_t \in \mathbb{Z}^k$? We can use recurrent neural networks for modelling $p(s_t | s_{t-1})$, e.g. Gated recurrent Units (GRU) or Long short-term memory models (LSTM).

